

# CSE 6220 Homework 2

Karl Hiner, @khiner6

## 1

Determine if the parallel prefix algorithm can be used to compute prefix sums of a sequence of  $n$  numbers based on the binary operation  $\oplus$  defined as:

The parallel prefix algorithm can be used to compute the prefix sums of a sequence of  $n$  numbers if and only if  $\oplus$  is a binary associative operator.

(a)  $a \oplus b = 2a + b$

Checking associativity,

$$\begin{aligned}(a \oplus b) \oplus c &\stackrel{?}{=} a \oplus (b \oplus c) \\ 2(2a + b) + c &\stackrel{?}{=} 2a + (2b + c) \\ 4a + 2b + c &\neq 2a + 2b + c\end{aligned}$$

This shows the binary operation defined as  $a \oplus b = 2a + b$  *is not* associative, and thus *cannot* be used by the parallel prefix algorithm to compute prefix sums.

(b)  $a \oplus b = \sqrt{a^2 + b^2}$

$$\begin{aligned}(a \oplus b) \oplus c &\stackrel{?}{=} a \oplus (b \oplus c) \\ \sqrt{\left(\sqrt{a^2 + b^2}\right)^2 + c^2} &\stackrel{?}{=} \sqrt{a^2 + \left(\sqrt{b^2 + c^2}\right)^2} \\ \left(\sqrt{a^2 + b^2}\right)^2 + c^2 &\stackrel{?}{=} a^2 + \left(\sqrt{b^2 + c^2}\right)^2 \\ a^2 + b^2 + c^2 &= a^2 + b^2 + c^2\end{aligned}$$

This shows the binary operation defined as  $a \oplus b = \sqrt{a^2 + b^2}$  *is* associative, and thus *can* be used by the parallel prefix algorithm to compute prefix sums.

## 2

In the game of Photosynthesis, points are given for trees that receive sunlight. Consider  $n$  trees  $T_0, T_1, \dots, T_{n-1}$  planted along a single row of spaces, and sunlight is colinear with this row of trees. The tree placement is modeled by an array  $A$  of size  $n$ , where  $A[i]$  denotes the height of

the tree  $T_i$ . A tree tall enough to get sunlight exposure scores photosynthesis points according to its height, i.e.,  $T_i$  is given  $A[i]$  points. However, a tree can also be blocked from sunlight by an earlier tree *of equal height or taller*, in which case the blocked tree receives no points.

Design a parallel algorithm to compute the total number of points for a configuration given by  $A$ , and compute its runtime.

Assume the array  $A$  is block-distributed across all nodes. Let  $p_k$  refer to the node holding the  $n/p$  elements  $A_k \equiv \left[ A\left[\frac{kn}{p}\right], \dots, A\left[\frac{(k+1)n}{p} - 1\right] \right]$ .

1. Perform parallel prefix with  $\oplus = \max$ , storing the result in array  $M_i$ . (Note that the max operation is binary associative, and thus can be used for parallel prefix.) After parallel prefix,  $M_i[j] \geq A[0, 1, \dots, \frac{ni}{p} + j]$ . Parallel prefix takes  $\Theta(\frac{n}{p} + \log p)$  time, with communication time  $\Theta((\tau + \mu) \log p)$ .
2. Use a right-shift to send the last element of  $M_{i-1}$  to its right neighbor processor  $i$ , storing this value in  $\hat{m}_i$ . This is necessary since each element of  $A_i$  will need to be compared with the *previous* max element. Communication cost:  $\Theta(\tau + \mu)$
3. Compute the local total tree score for each node, as follows:
  - Initialize the running total local score  $s_k \leftarrow 0$ .
  - For each tree height in the local list  $A_i$ , if  $A_i[j] > M_i[j - 1]$ , add the tree height to the total local score,  $s_k += A_j[j]$ . For  $j - 1 = -1$ , use  $\hat{m}_i$  instead of (nonexistent)  $M_i[-1]$ .

This step takes  $\Theta(\frac{n}{p})$  time.

4. Use the parallel sum algorithm to sum all local tree scores to find the total tree score,  $S = \sum_k s_k$ . This final step takes  $\Theta(\log p)$  time.

Computation time:  $\Theta(2(\frac{n}{p} + \log p))$ .

Communication time:  $\Theta((\tau + \mu)(\log p + 1))$ .

### 3

A sequence of nested parenthesis is said to be well-formed if

1. there are an equal number of left and right parenthesis, and
2. each right parenthesis is matched by a left parenthesis that occurs to its left in the sequence.

For example,  $(( ( ) ( ) ) ( ) )$  is well-formed, but  $( ) ) ($  is not.

There is a nested parenthesis sequence of length  $n$  distributed across  $p$  processors using block decomposition. Design a parallel algorithm to determine if it is well-formed and specify its run-time.

1. Assign  $-1$  to a closed parenthesis, and  $1$  to an open parenthesis. Then each processor  $i$  has a list of values  $A_i \in (-1, 1)^{\frac{n}{p}}$ .
2. Use parallel prefix with  $a \oplus b = a + b$  (parallel prefix sum).
3. If any (global) prefix sum element is negative, return "not well-formed".

4. If all prefix sums are non-negative, and the final prefix sum (the total sum) is equal to 0, the algorithm returns "well-formed". Otherwise, return "not well-formed".

The runtime of this algorithm is the same as parallel prefix sum,  $\Theta(\frac{n}{p} + \log p)$ . Unlike prefix sum, however, the best-case runtime is  $\Theta(1)$ , since the first processor could return "not well-formed" if its first element is a closed parenthesis  $(-1)$ .

## 4

Let  $A$  be an array of  $n$  elements and  $L$  be a boolean array of the same size. We want to assign a unique rank in the range  $1, 2, \dots, n$  to each element of  $A$  such that for any  $i < j$ :

- If  $L[i] = L[j]$ ,  $A[i]$  has lower rank than  $A[j]$ .
- If  $L[i] = 0$  and  $L[j] = 1$ ,  $A[i]$  has lower rank than  $A[j]$ .
- If  $L[i] = 1$  and  $L[j] = 0$ ,  $A[j]$  has lower rank than  $A[i]$ .

Design a parallel algorithm to compute the ranks and specify its run-time. (Hint: Think of  $L$  as specifying labels. Then, all elements with 0 label receive lower ranks than any element with label 1. Within the same label, ranks are given in left to right order as per array  $A$ .)

Assume arrays  $A$  and  $L$  are block-distributed across  $p$  processors, such that each processor holds  $\frac{n}{p}$  contiguous elements from each array in local arrays  $A_i$  and  $L_i$ .

1. Compute the number of 0s and 1s in each node's  $L_i$  array, and store in  $\mathbf{n}_i = \begin{bmatrix} \text{num 0s} \\ \text{num 1s} \end{bmatrix}$ .  
This step takes  $\Theta(\frac{n}{p})$  time.
2. Perform a parallel prefix operation across  $\mathbf{n}_i$ , using vector sum as the operator, and store the results in  $\mathbf{N}_i = \sum_{j \leq i} \mathbf{n}_j$ . (Note that after this step,  $\mathbf{N}_i - \mathbf{n}_i$  gives the total number of zeros/ones *before* node  $i$ .) This step takes  $\Theta(2 \log p)$  computation time and  $\Theta((\tau + 2\mu) \log p)$  communication time.
3. Broadcast the total number of 0s in  $L$  (stored in  $\mathbf{N}_p[0]$ , where  $p$  is the index of the last processor), to every other node. Store the result as  $Z$  in each node. This takes  $\Theta((\tau + \mu) \log p)$  communication time.
4. For each node, traverse its  $L_i$  and assign ranks to elements in  $A$  as follows, storing the result in local rank array  $R_i$ :

$$R_i[j] = \begin{cases} c_0 + \mathbf{N}_i[0] - \mathbf{n}_i[0], & L_i[j] = 0 \quad (c_0 + \text{num 0s before me}) \\ c_1 + Z + \mathbf{N}_i[1] - \mathbf{n}_i[1], & L_i[j] = 1 \quad (c_1 + \text{total 0s} + \text{num 1s before me}) \end{cases},$$

where  $c_0/c_1$  are the running count of 0s/1s encountered so far during the traversal. This step takes  $\Theta(\frac{n}{p})$  time.

Computation time:  $\Theta(2(\frac{n}{p} + \log p))$ .

Communication time:  $\Theta((2\tau + 3\mu) \log p)$ .

*Note:* Instead of using  $Z$ , the total number of 0s in  $L$ , as a rank offset when  $L_i[j] = 1$ , we could use any  $N \geq Z$ . For example, we could use  $n$ . We would still need to broadcast  $n$ , as we did for  $Z$ , but it could be done in the first step. Alternatively, if we were allowed to provide  $n$  to each processor along with its  $L_i$  and  $A_i$  (as in PA1), we could skip the broadcast step entirely, resulting in a total communication time of  $\Theta((\tau + 2\mu) \log p)$ .

## 5

*Invent Segmented Parallel Prefix:* Segmented parallel prefix is a generalization of the parallel prefix problem where the prefix sums need to be restarted at specified positions. Consider array  $X$  containing  $n$  numbers and a boolean array  $B$  of the same size. We wish to compute prefix sums on  $X$  but the sum resets at every position  $i$  where  $B[i] = 1$ . Formally, we wish to compute array  $S$  of size  $n$  such that

$$S[0] = X[0]$$

$$S[i] = \begin{cases} S[i-1] + X[i], & \text{if } B[i] = 0 \\ X[i], & \text{if } B[i] = 1 \end{cases}$$

Design parallel segmented prefix algorithm and specify its run-time.

(Hint: The problem can be transformed into a standard prefix sums problem.)

We are told this can be transformed into a standard prefix sums problem. So we must find a binary associative operator that satisfies the given definition for  $S$ .

Simplifying the  $S[i]$  case:

$$\begin{aligned} S[i] &= \begin{cases} S[i-1] + X[i], & B[i] = 0 \\ X[i], & B[i] = 1 \end{cases} \\ &= X[i] + \begin{cases} S[i-1], & B[i] = 0 \\ 0, & B[i] = 1 \end{cases} \\ &= X[i] + S[i-1](1 - B[i]) \end{aligned}$$

For the case of  $i = 1$ , we can solve this using the base case of  $S[0] = X[0]$  to get rid of the reference to the previous  $S$  element on the right:

$$\begin{aligned} S[1] &= X[1] + S[0](1 - B[1]) \\ &= X[1] + X[0](1 - B[1]) \end{aligned}$$

Our binary operator must act on element types containing all the information needed in this expression. Let this element type be defined as  $\mathbf{A}_i \equiv \begin{bmatrix} X_i \\ B_i \end{bmatrix} \equiv \begin{bmatrix} X[i] \\ B[i] \end{bmatrix}$ . That is,  $\mathbf{A}_i$  is the two-element vector formed by combining the  $i$ th element of  $X$  with the  $i$ th element of  $B$ . Then we can consider the inputs  $X$  and  $B$  as being provided in a single  $2 \times n$  matrix  $\mathbf{A}$ , whose columns  $\mathbf{A}_i$  we block-distribute across  $p$  processes. (We do not actually need to create these elements - this is only a conceptual reformulation with no incurred cost.)

From our solution to  $S[1]$  above, we know exactly how to compute the first item (the numeric value) in our binary associative operator  $\oplus$ , given two consecutive columns of  $\mathbf{A}$  (since  $S[1] \equiv (\mathbf{A}_0 \oplus \mathbf{A}_1)_0$ ). We then only need to consider how to compute the second item (the boolean value). I must admit I did this step by trial and error, so I don't provide a derivation here, but found that the logical-or (max) works.

Thus, our boolean operator is:

$$\mathbf{A}_i \oplus \mathbf{A}_j \implies \begin{bmatrix} X_i \\ B_i \end{bmatrix} \oplus \begin{bmatrix} X_j \\ B_j \end{bmatrix} = \begin{bmatrix} X_j + X_i(1 - B_j) \\ \max(B_i, B_j) \end{bmatrix}$$

Checking associativity:

$$\begin{aligned} (\mathbf{A}_a \oplus \mathbf{A}_b) \oplus \mathbf{A}_c &\stackrel{?}{=} \mathbf{A}_a \oplus (\mathbf{A}_b \oplus \mathbf{A}_c) \\ \left( \begin{bmatrix} X_a \\ B_a \end{bmatrix} \oplus \begin{bmatrix} X_b \\ B_b \end{bmatrix} \right) \oplus \begin{bmatrix} X_c \\ B_c \end{bmatrix} &\stackrel{?}{=} \begin{bmatrix} X_a \\ B_a \end{bmatrix} \oplus \left( \begin{bmatrix} X_b \\ B_b \end{bmatrix} \oplus \begin{bmatrix} X_c \\ B_c \end{bmatrix} \right) \\ \left( \begin{bmatrix} X_b + X_a(1 - B_b) \\ \max(B_a, B_b) \end{bmatrix} \right) \oplus \begin{bmatrix} X_c \\ B_c \end{bmatrix} &\stackrel{?}{=} \begin{bmatrix} X_a \\ B_a \end{bmatrix} \oplus \left( \begin{bmatrix} X_c + X_b(1 - B_c) \\ \max(B_b, B_c) \end{bmatrix} \right) \\ \begin{bmatrix} X_c + (X_b + X_a(1 - B_b))(1 - B_c) \\ \max(\max(B_a, B_b), B_c) \end{bmatrix} &\stackrel{?}{=} \begin{bmatrix} X_c + X_b(1 - B_c) + X_a(1 - \max(B_b, B_c)) \\ \max(B_a, \max(B_b, B_c)) \end{bmatrix} \\ \begin{bmatrix} X_c + X_b(1 - B_c) + X_a(1 - B_b)(1 - B_c) \\ \max(B_a, B_b, B_c) \end{bmatrix} &\stackrel{?}{=} \begin{bmatrix} X_c + X_b(1 - B_c) + X_a(1 - \max(B_b, B_c)) \\ \max(B_a, B_b, B_c) \end{bmatrix} \end{aligned}$$

Cancelling out all like terms, we are left with

$$\begin{aligned} (1 - B_b)(1 - B_c) &\stackrel{?}{=} 1 - \max(B_b, B_c). \\ B_b + B_c - B_b B_c &= \max(B_b, B_c). \end{aligned}$$

This equality can be readily verified with a truth table (not provided here).

Finally, here is the algorithm, assuming  $X$  and  $B$  are block distributed.

1. Perform parallel prefix using  $\begin{bmatrix} X_i \\ B_i \end{bmatrix} \oplus \begin{bmatrix} X_j \\ B_j \end{bmatrix} = \begin{bmatrix} X_j + X_i(1 - B_j) \\ \max(B_i, B_j) \end{bmatrix}$ , storing the resulting two-element vectors as columns in the  $\left(2 \times \frac{n}{p}\right)$  matrix  $\mathbf{A}_i$ .
2. The array  $S$  can be read directly from the first row of each  $\mathbf{A}_i$ .

Computation time:  $\Theta\left(4\left(\frac{n}{p} + \log p\right)\right)$  (parallel prefix with  $\approx 4$  flops per op).

Communication time:  $\Theta((\tau + 2\mu) \log p)$  (parallel prefix with 2D elements).